

Scaling and monitoring your puppetserver for thousands of clients


Includes all pitfalls!

\$ whoami

- Tim 'bastelfreak' Meusel
- DevOps Engineer at GoDaddy EMEA
- Puppet Contributor since 2012
- Merging stuff on Vox Pupuli since 2015
- Vox Pupuli PMC member

starting
point

Puppetserver




@bastelsblog for @voxpupuliorg

starting
point

Puppetserver

PuppetDB




@bastelsblog for @voxpupuliorg

starting
point

Puppetserver

PuppetDB

Postgresql



@bastelsblog for @voxpupuliorg


starting
point

Puppetserver

PuppetDB

Postgresql

PuppetBoard



starting
point


Puppetserver

PuppetDB

Postgresql

PuppetBoard

Foreman



@bastelsblog for @voxpupuliorg

starting
point

Puppetserver


PuppetDB

Postgresql

PuppetBoard

Foreman

gnatsd

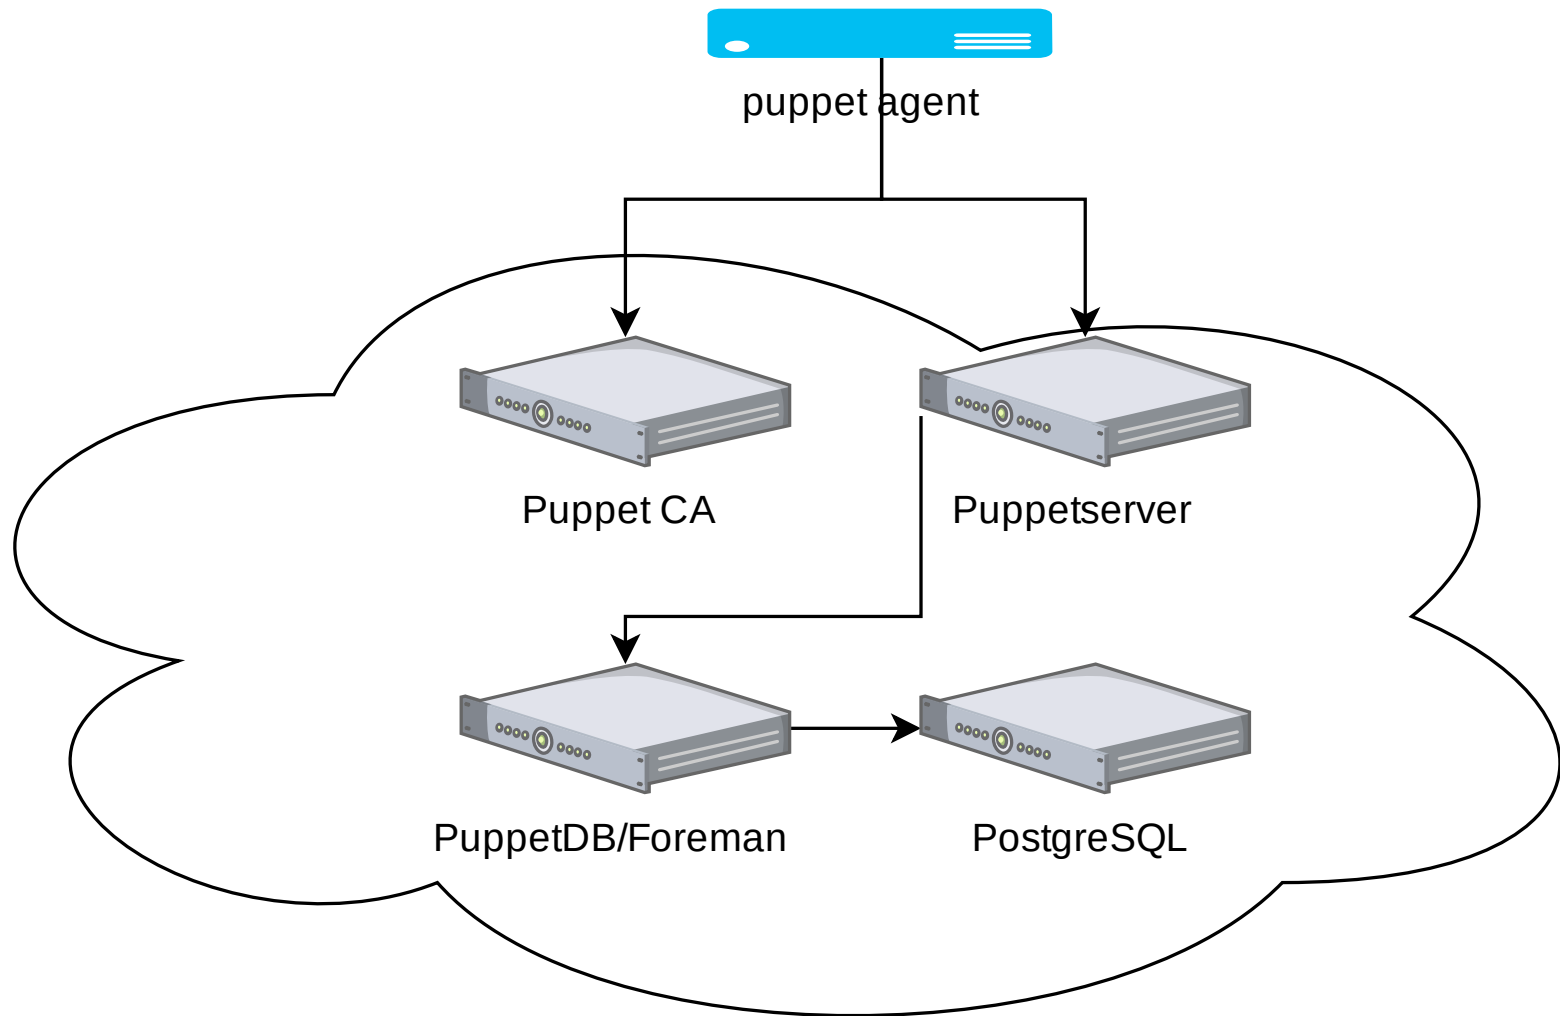


@bastelsblog for @voxpupuliorg

Coworker: Puppet agents are slow, can you please check the monitoring?

Coworker: Puppet agents are slow, can you please check the monitoring?

Uhm, which monitoring?



This is the cloud. The cloud is a half rack in your favourite datacenter

Improvements

Puppetserver

- Uses jruby within a JVM
- One jruby instance compiles one catalog at a time
- More instances => more catalogs per minute
- Use [theforeman/puppet](#) to configure puppetserver

```
---  
puppet::server_max_active_instances: %{facts.processors.count}
```

Improvements

Puppetserver

- Puppetserver can cache code by loading it from disk to ram
- + Minimal decreased compilation time for each catalog
- - You should clean the cache after each deploy

```
---  
puppet::server_environment_class_cache_enabled: true
```

- cleaning the cache:

```
print hostcert) \  
stprivkey) \  
cacert) \  
  
server):8140/puppet-admin-api/v1/environment-cache?environment=production
```

Improvements

Puppetserver

- `theforeman/puppet` creates the `development` and `production` environment
- `r10k` purges unknown environments
- `theforeman/puppet` restarts puppetserver if it creates an environment
- I have no git branches named `development` nor `production` in my control repo...
- Each environment deploy lead to a restarted puppetserver for weeks

```
---  
# don't create development/production env  
puppet::server_environments: []  
# don't create /etc/puppetlabs/code/environments/common  
puppet::server_common_modules_path: ''
```

Improvements

Puppetserver

- JVM has a configureable minimal and maximal amount of memory to allocate
- Memory is shared across all jruby instances (and other threads)
- Puppet docs suggest that minimal=maximal memory
 - That is based on Java 6 docs, so probably outdated
- Required memory per instance depends entirely on the codebase (modules)
- 2GB seem to work out fine for my setup

```
# How do I do this with hiera?  
$cpu_count_twice = $facts['processors']['count'] * 2  
$cpu_count = $facts['processors']['count'] * 1  
class{'puppet':  
  server_jvm_min_heap_size => "${cpu_count}G",  
  server_jvm_max_heap_size => "${cpu_count_twice}G",  
}
```


Improvements

Puppetserver

Foreman

- foreman supports caching out of the box
- we use `saz/memcached` to provision `memcached`

```
# 50GB of cache
memcached::max_memory: 51200
foreman::plugin::memcache::hosts:
  - 127.0.0.1
```

```
include memcached
include foreman::plugin::memcache
```

Improvements

Puppetserver

Foreman

- `passenger-status` says it only runs with a single process..
- `theforeman/foreman` uses `puppetlabs/apache` to configure `passenger`

```
--  
apache::mod::passenger::passenger_max_pool_size: %{facts.processors.count}  
apache::mod::passenger::passenger_min_instances: %{facts.processors.count}
```

Improvements

Puppetserver

Foreman

- We use **theforeman/foreman** to manage foreman
- Tuned puppetserver results in more requests to foreman

```
---  
# default is 5  
foreman::db_pool: 20  
foreman::keepalive: true  
foreman::max_keepalive_requests: 1000  
foreman::keepalive_timeout: 180
```

Improvements

Puppetserver

Foreman

PostgreSQL

- This deserves a dedicated conference (there actually is)
- Attend the [PostgresConf](#) or the PostgreSQL DevRoom at [FOSDEM](#)
- Ask people in #postgresql on freenode
- Use at least postgres 10 and rely on the upstream repos if possible
- Don't use harddrive, SATA/SAS SSDs and NVMe SSDs are the way to go
- Execute pg tune

```
pgtune -i /var/lib/pgsql/10/data/postgresql.conf -o postgresql5.conf
```

```
postgresql::server::config_entry{'max_connections':  
  value => 400,  
}
```

Improvements

- Simply JVM service that scales good with the amount of threads
- It's a RESTful service that stores data in PostgreSQL

Puppetserver

Foreman

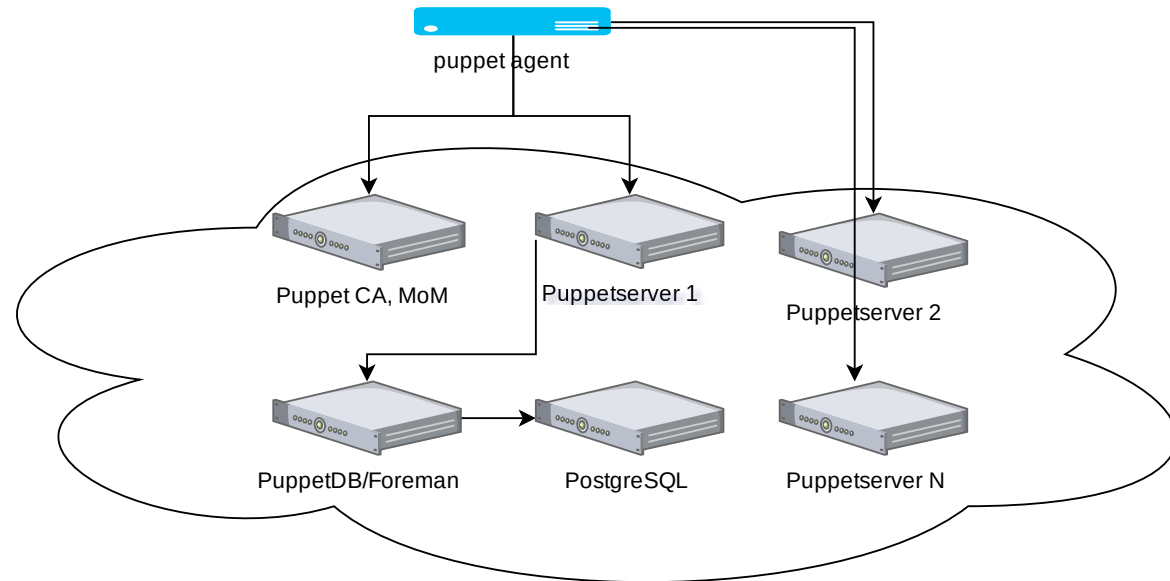
PostgreSQL

PuppetDB

```
puppetdb::server::java_args:  
  '-Xmx': '8192m'  
  '-Xms': '2048m'  
puppetdb::server::node_ttl: '14d',  
puppetdb::server::node_purge_ttl: '14d',  
puppetdb::server::report_ttl: '999d'  
# default is 50  
puppetdb::server::max_threads: 100  
# default is processorcount / 2  
puppetdb::server::command_threads: %{facts.processors.count}  
# default is 4, have your database in mind  
puppetdb::server::concurrent_writes: 8  
puppetdb::server::automatic_dlo_cleanup: true
```

Scaling

The Idea



This is the cloud. The cloud is one rack in your favourite datacenter

Scaling

The Idea

tk-jetty9


Puppetserver docs:

- **selector-threads**: This sets the number of selectors that the webserver will dedicate to processing events on connected sockets for unencrypted HTTPS traffic. No known upper limit
- **ssl-selector-threads**: same as **selector-threads**, just for HTTPS. "Defaults to the number of virtual cores on the host divided by 2, with a minimum of 1 and maximum of 4"

Scaling

The Idea

tk-jetty9



Scaling

The Idea

tk-jetty9

The Setup

- Deploy nginx on each Puppetserver server to terminate TLS
 - Increase default threadcount from 1 to \$more...
- Bind puppetserver to localhost with http
 - And don't require TLS client certificates
- Setup consul for dynamic loadbalancing across all puppetservers
- Soonish available at github.com/bastelfreak/puppetcontrolrepo

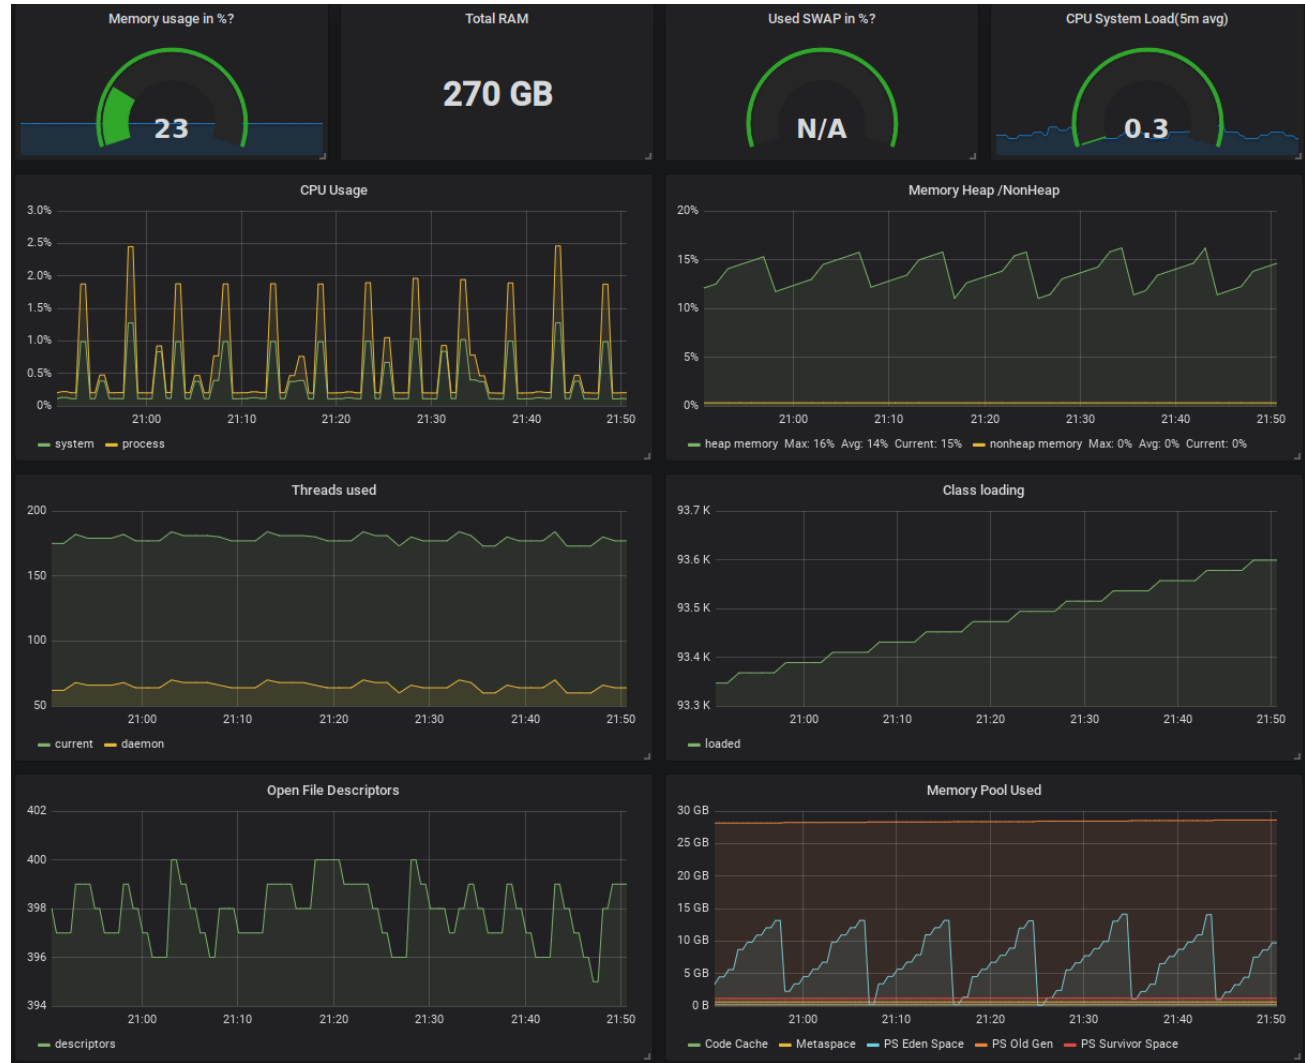
Metrics

- Puppetserver can expose graphite
 - A graphite stack looks very complicated
- Puppetserver has a metrics API and exposes **JMX data** via **Jolokia**
- We can load a prometheus exporter into the JVM to write metrics into a prometheus instance

```
puppet::server_jvm_extra_args:  
  - '-javaagent:prometheus.jar=127.0.0.1:9020:config.yaml'
```

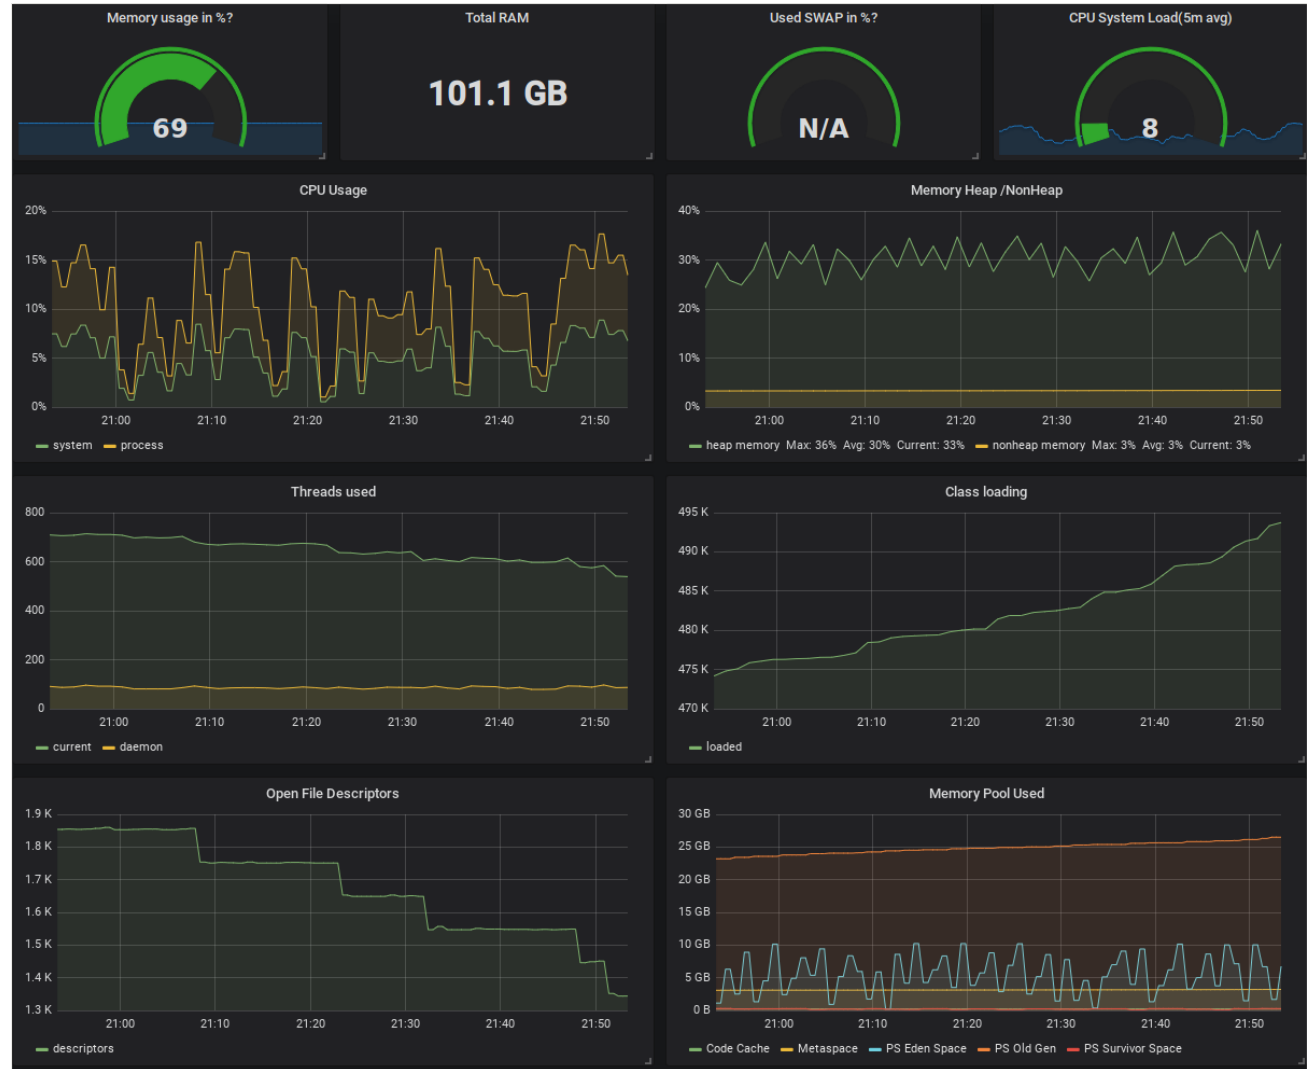
Metrics

Generic JMX



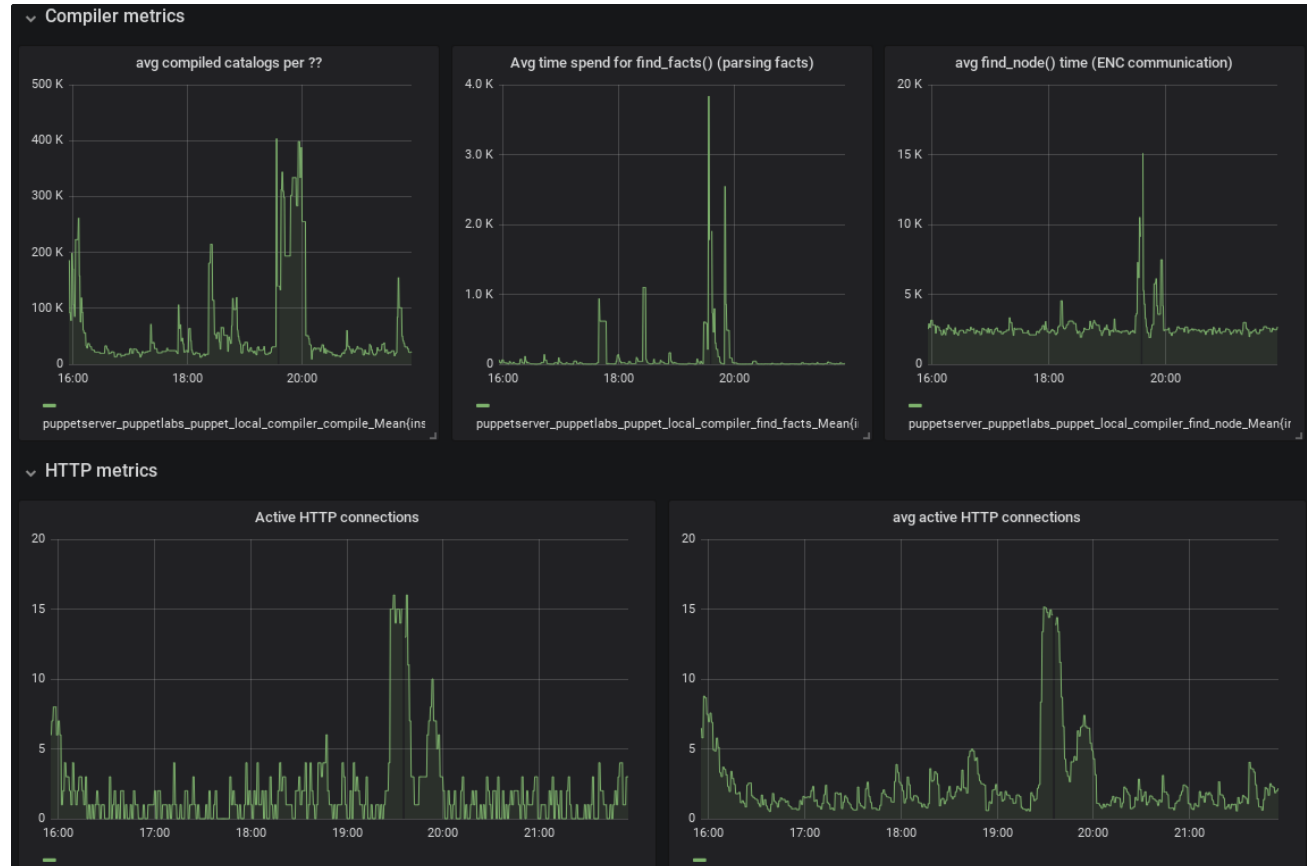
Metrics

Generic JMX




Metrics

Puppetserver



Metrics

Puppetserver



Summary

Scaling a Puppetserver stack

- It's a complex distributed system with many tunables and pitfalls
- Start with proper monitoring instead of guessing
- Best practice [controlrepo](#) with all tuneables, explanations, unit/acceptance tests
- Contact: tim@bastelfreak.de or bastelfreak on freenode
- [Collection of related talks](#)

Thanks for your attention!